

Introduction to Parallel Computing

Byoung-Do (BD) Kim, PhD
Associate Chief Research Information Officer
Director, Center for Advanced Research Computing

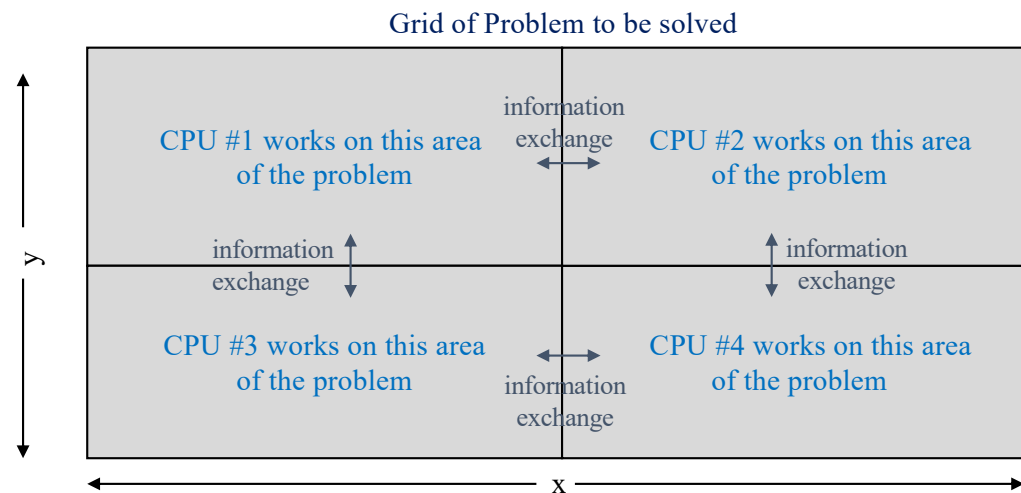


Outline

1. Introduction to Parallel Computing
2. Theoretical background
3. Types of parallel computing systems
4. Programming models
5. Examples
6. Hands-on session

What is Parallel Computing?

- Parallel computing: use of multiple processors or computers working together on a common task.
 - Each processor works on its section of the problem
 - Processors can exchange information





Why Parallel Computing?

- Moore's Law?
 - Processor speed is no longer double every 18-24 months
 - Multi-core is the norm
- Parallel computing allows one to:
 - solve problems that don't fit on a single CPU
 - solve problems that can't be solved in a reasonable time
- We can solve...
 - larger problems
 - faster
 - more cases in a given time

Supercomputing Applications

- Black Hole Simulation
 - [Black Hole simulation](#)
 - [NASA article](#)
- Formular 1 Racing Car Aerodynamics:
 - [F1 aerodynamics](#)
 - [Flex wing debate](#)
- Supercomputing vs. COVID-19
 - [Drug discovery for COVID-19 using supercomputer](#)
- AI & Supercomputer
 - [DeepMind AlphaGo defeated professional Go player 4:1](#)
 - [AlphaGo Movie](#)
- Stock Market High-Frequency Trading



Advanced Research Computing
Enabling scientific breakthroughs at scale

Top 500

- List of fastest supercomputers in the world
 - [Top500](#)
 - [Green500](#)
 - [List statistics](#)

IBM-built Summit supercomputer at
Oak Ridge National Laboratory

- 148.6 Petaflops (200PF peak)
- 9,216 Power9 CPUs (203K cores)
- 27,648 Nvidia V100 GPUs
- 13MW



Advanced Research Computing
Enabling scientific breakthroughs at scale

Limits of Parallel Computing

- Theoretical Upper Limits
 - Amdahl's Law
- Practical Limits
 - Load balancing
 - Non-computational sections
 - Communication overhead
- Other Considerations
 - time to re-write code

Parallel Computing: Theoretical background

Theoretical Upper Limit to Performance

- All parallel programs contain:
 - parallel sections (we hope!)
 - serial sections (unfortunately)
- Serial sections limit the parallel effectiveness
- Amdahl's Law states this formally

Amdahl's Law


- Amdahl's Law places a strict limit on the speedup that can be realized by using multiple processors.

- Speedup: $S = \frac{T_{serial}}{T_{parallel}}$

- Effect of multiple processors on speed up: $S = \frac{1}{f_s + \frac{f_p}{N}}$

where

- f_s = serial fraction of code
- f_p = parallel fraction of code
- N = number of processors




- Amdahl's law in multi-core era: <https://research.cs.wisc.edu/multifacet/amdahl/>



Illustration of Amdahl's Law

- It takes only a small fraction of serial content in a code to degrade the parallel performance.




From Wikipedia: Amdahl's Law
https://en.wikipedia.org/wiki/Amdahl%27s_law



Advanced Research Computing
Enabling scientific breakthroughs at scale

Practical Limit: Amdahl's Law vs. Reality

- Amdahl's Law provides a theoretical upper limit on parallel speedup assuming that there are no costs for *communications*. In reality, communications will result in a further degradation of performance.

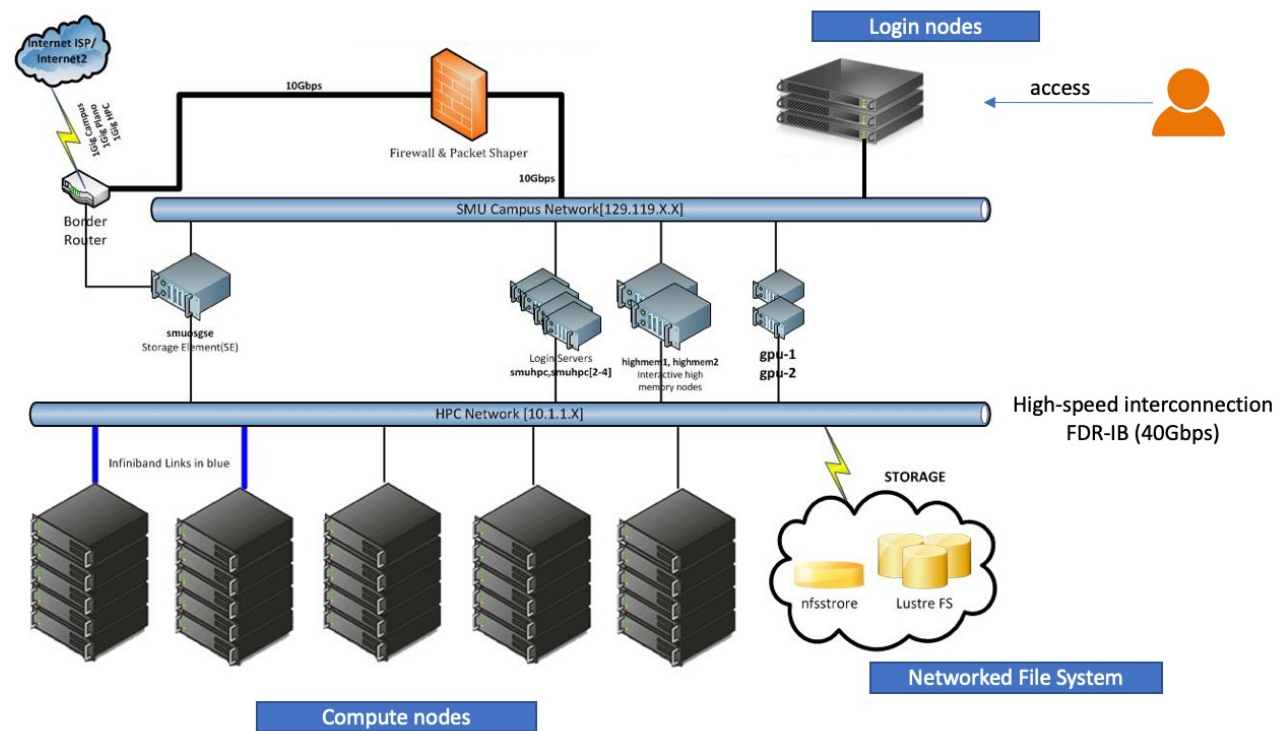


Practical Limit: Amdahl's Law vs. Reality

- In reality, the situation is even worse than predicted by Amdahl's Law
 - Scheduling (shared processors or memory)
 - Communications
 - I/O
- Writing effective parallel applications is difficult!
 - Load balance is important
 - Communication can limit parallel efficiency
 - Serial time can dominate
- Is it worth your time to rewrite your application?
 - Do the CPU requirements justify parallelization?
 - Will the code be used just once?

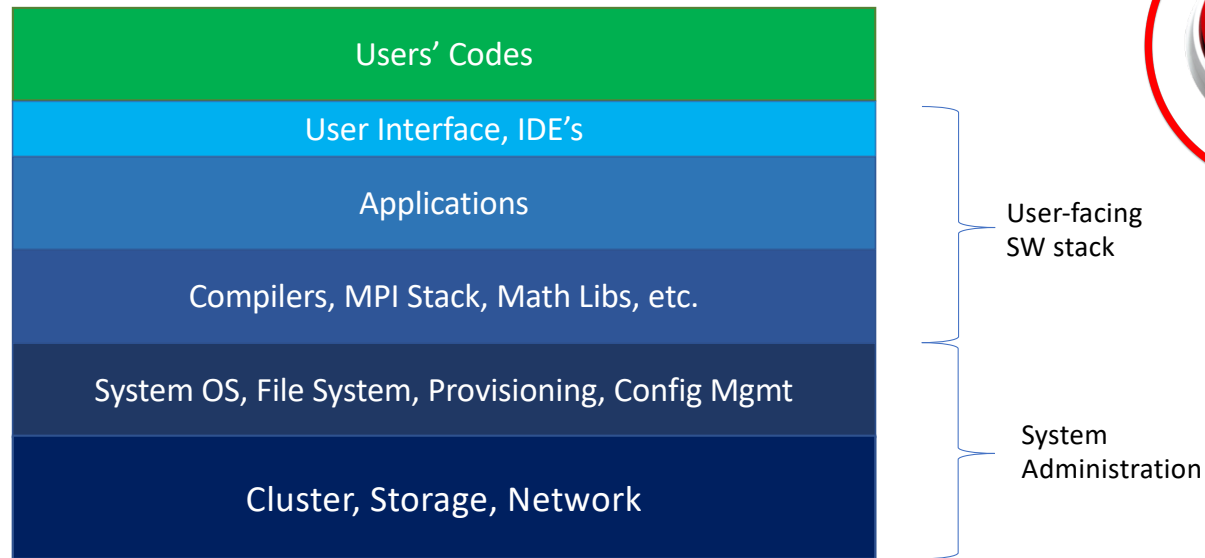
Parallel System Architecture

HPC in a Nutshell: Cluster System Architecture




Advanced Research Computing
Enabling scientific breakthroughs at scale

HPC in a Nutshell: **SW Stack**




Shared vs. Distributed Memory



Shared memory:


- Single address space.
- All processors have access to a pool of shared memory.
- Methods of memory access :
 - Bus, Crossbar
- Programming model: OpenMP



Distributed memory:

- Each processor has its own local memory.
- Must do message passing to exchange data between processors.
- Methods of memory access :
 - Various topological interconnection
- Programming model: MPI

Multicore with Accelerators



- A limited number of processors N have access to a common pool of shared memory
- To use more than N processors requires data exchange over a network
- Communication details increasingly complex
 - Cache access vs. Main memory access
 - Quick Path / Hyper Transport socket connections
 - Node to node connection via high-speed network
- Load balancing critical for performance
- Requires specific libraries and compilers (CUDA, OpenCL, ACC, etc.)
- Huge performance boost is possible



Parallel Programming Models

Parallel Programming Models

- Data Parallelism
 - Each processor performs the same task on different data
- Task Parallelism
 - Each processor performs a different task on the same data (or on different data)
- Most applications fall between these two

Data Parallel Programming Example

- One code will run on 2 CPUs
- Program has array of data to be operated on 2 CPUs, array is split into two parts.

```
program:
...
if CPU=a then
  low_limit=1
  upper_limit=50
elseif CPU=b then
  low_limit=51
  upper_limit=100
end if
do I = low_limit,
  upper_limit
  work on A(I)
end do
...
end program
```

CPU A

```
program:
...
low_limit=1
upper_limit=50
do I= low_limit,
  upper_limit
  work on A(I)
end do
...
end program
```

CPU B

```
program:
...
low_limit=51
upper_limit=100
do I= low_limit,
  upper_limit
  work on A(I)
end do
...
end program
```



Task Parallel Programming Example

- One code will run on 2 CPUs
- Program has 2 tasks (a and b) to be done by 2 CPUs

```
program:  
...  
initialize  
...  
if CPU=a then  
    do task a  
elseif CPU=b then  
    do task b  
end if  
...  
end program
```

CPU A

```
program:  
...  
initialize  
...  
do task a  
...  
end program
```

CPU B


```
program:  
...  
initialize  
...  
do task b  
...  
end program
```



Single Program Multiple Data

- SPMD: dominant programming model for shared and distributed memory machines.
 - One source code is written
 - Code can have conditional execution based on which processor is executing the copy
 - All copies of code start simultaneously and communicate and sync with each other periodically

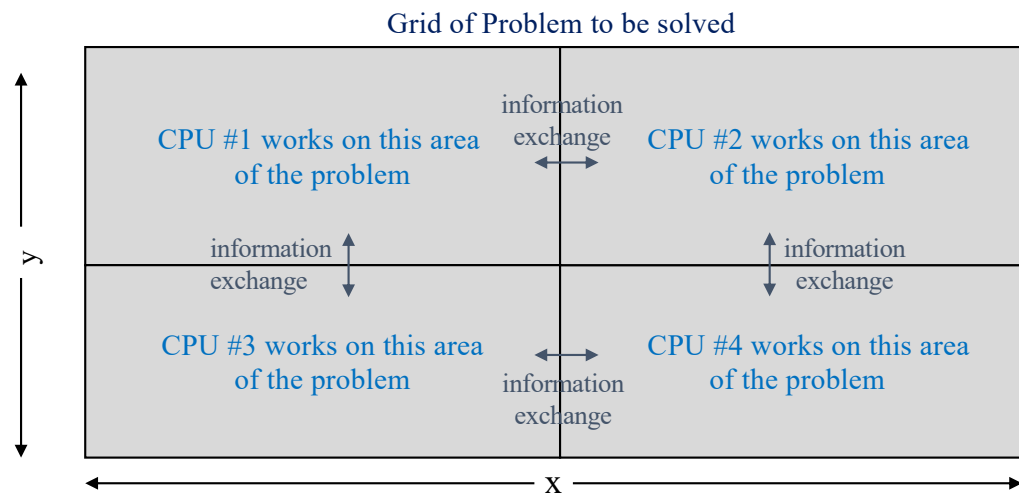
SPMD Model



- Ideal programming model in multi-node system environment

Data Decomposition

- For distributed memory systems, the 'whole' grid or sum of particles is decomposed to the individual processors
 - Each CPU works on its section of the problem
 - CPUs/Nodes can exchange information



Domain Decomposition Example


- EX) 2-D wave propagation problem

Starting partial
differential equation:

$$\frac{\partial \Psi}{\partial t} = D \cdot \frac{\partial^2 \Psi}{\partial x^2} + B \cdot \frac{\partial^2 \Psi}{\partial y^2}$$

Finite Difference
Approximation:

$$\frac{f_{i,j}^{n+1} - f_{i,j}^n}{\Delta t} = D \cdot \frac{f_{i+1,j}^n - 2f_{i,j}^n + f_{i-1,j}^n}{\Delta x^2} + B \cdot \frac{f_{i,j+1}^n - 2f_{i,j}^n + f_{i,j-1}^n}{\Delta y^2}$$




Advanced Research Computing
Enabling scientific breakthroughs at scale

Numerical discretization of PDE

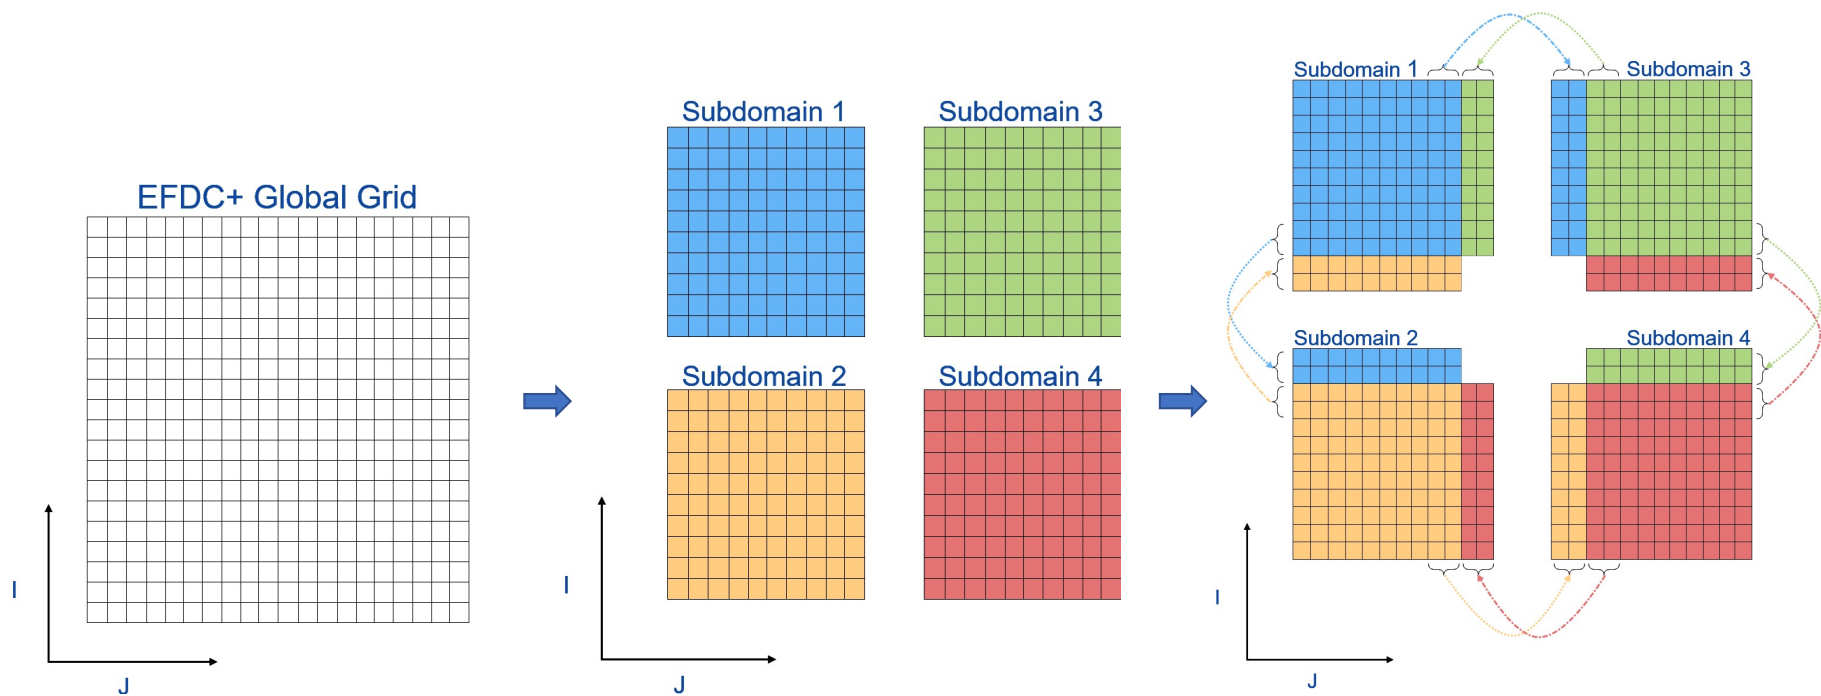
$$\left(\frac{\partial^2 u}{\partial x^2}\right) \approx \frac{u(x-\Delta x, y) - 2u(x, y) + u(x+\Delta x, y)}{(\Delta x)^2},$$

$$\left(\frac{\partial^2 u}{\partial y^2}\right) \approx \frac{u(x, y-\Delta y) - 2u(x, y) + u(x, y+\Delta y)}{(\Delta y)^2}$$



Implementation of Domain Decomposition

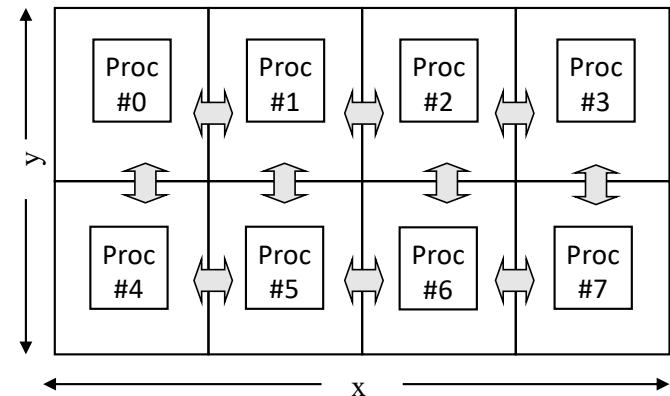
<https://www.eemodelingsystem.com/efdc-insider-blog/domain-decomposition-details>, by Nghiem Tien Lam




Advanced Research Computing
Enabling scientific breakthroughs at scale

Information Propagation between Processes

- How do we decide on what & how much of information should be passed along between decomposed computational domain?
- based on PDE's *characteristics*




Elliptic PDE



* Information propagation speed:


Undefined

Parabolic PDE



Infinite

Hyperbolic PDE



Finite



Advanced Research Computing
Enabling scientific breakthroughs at scale


MPI: Message Passing Interface

MPI: Message Passing Interface

- MPI 1 was released in 1994, MPI 2.0 in 1997, and MPI 3.0 in 2012
- Distributed memory programming
- Ideal for multi-node parallelization
- Can use with OpenMP for better scalability
- Distributed memory systems have separate address spaces for each processor
 - Local memory accessed faster than remote memory
 - Data must be manually decomposed
 - MPI is the standard for distributed memory programming



Message Passing Communication



- Processes in message passing program communicate by passing messages
- Basic message passing primitives
 - `MPI_Send` (parameters list)
 - `MPI_Receive` (parameter list)
- These calls are **blocking**: the source processor issuing the send/receive cannot move to the next statement until the target processor issues the matching receive/send.

Communicators

- MPI uses **MPI_Comm** objects to define subsets of processors which may communicate with one another.
- Two common functions for interacting with an **MPI_Comm** object are:
- **MPI_Comm_size(MPI_Comm_World, int *np)**
 - Gets the number of processes in a run, *NP*
- **MPI_Comm_rank(MPI_Comm_World, int *rank)**
 - Gets the rank of the current process
 - returns a value between 0 and *NP*-1 inclusive
- Both are typically called just after **MPI_Init**.

Sample MPI code (C)

```
#include <mpi.h>
[other includes]

int main(int argc, char **argv){
    int ierr, np, rank;
    [other declarations]

    ierr = MPI_Init(&argc, &argv);
    ierr = MPI_Comm_size(MPI_COMM_WORLD, &np);
    ierr = MPI_Comm_rank(MPI_COMM_WORLD, &rank);
        :
    [actual work goes here]
        :
    MPI_Finalize();
}
```



Sample MPI code (F90)

```
program sample-mpi
  use mpi
  [other includes]

  integer :: ierr, np, rank
  [other declarations]

  call mpi_init(ierr)
  call mpi_comm_size(MPI_COMM_WORLD, np, ierr)
  call mpi_comm_rank(MPI_COMM_WORLD, rank, ierr)
  :
  [actual work goes here]
  :
  call mpi_finalize(ierr)
end program
```




Point-to-Point Communication

- Sending data from one point (process/task) to another point (process/task)
- One task sends while another receives

```
MPI_Send(buf, count, datatype, dest, tag, comm);
```

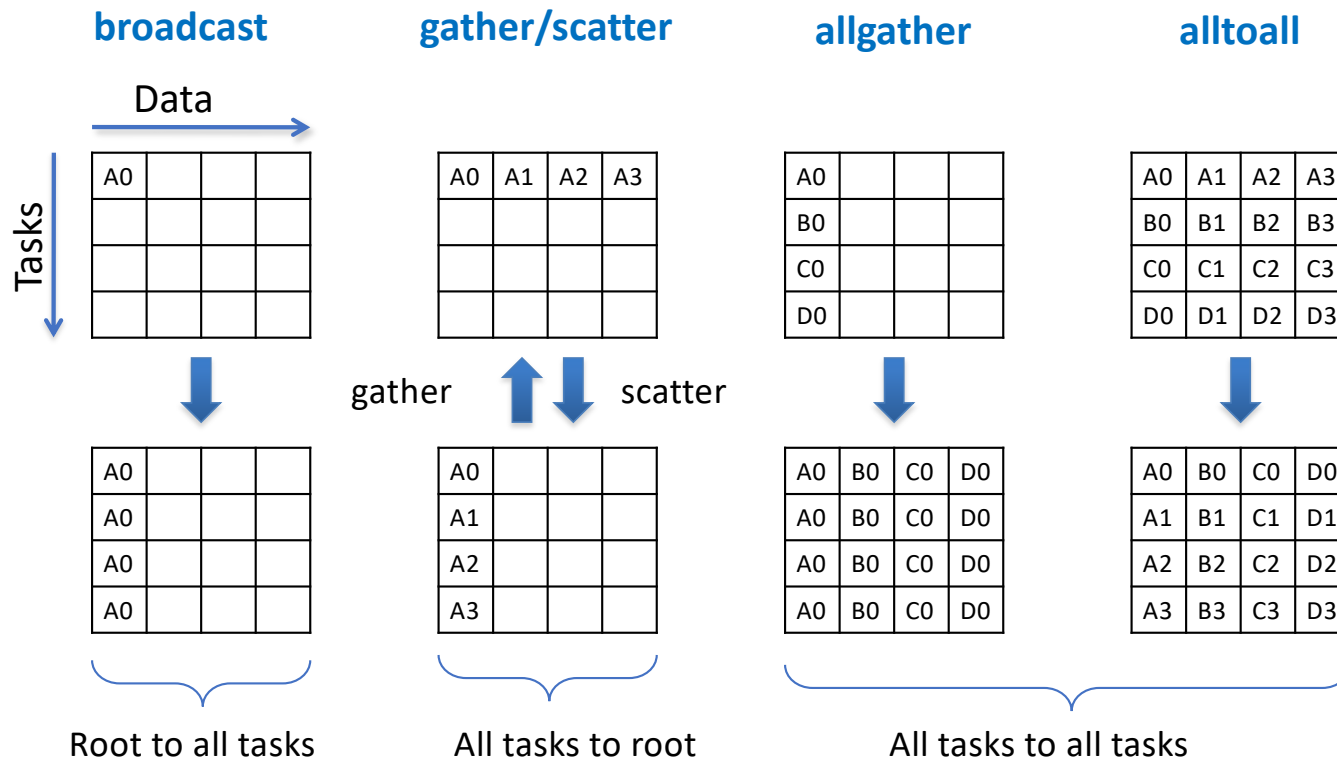
```
MPI_Recv(buf, count, datatype, source, tag, comm, status);
```



Collective Communication

- Defined as communication between > 2 processors
 - One-to-many
 - Many-to-one
 - Many-to-many
- A collective operation requires that all processes within the communicator group call the *same* collective communication function with matching arguments.
- Type of collective operations
 - **Synchronization** (MPI_Barrier)
 - **Data Movement** (MPI_Bcast/Scatter/Gather/Allgather/AlltoAll)
 - **Computation** (MPI_Reduce/Allreduce/Scan)

Collective Communication Visualization



MPI Programming

MPI Programming: Basic Structure

Every MPI program needs these:

```
#include <mpi.h> /* the mpi include file */
int main(int argc, char *argv[])
{
    /* Initialize MPI */
    ierr = MPI_Init(&argc, &argv);
    /* How many total processor are there */
    ierr = MPI_Comm_size(MPI_COMM_WORLD, &nprocs);
    /* What process am I (what is my rank? */
    ierr = MPI_Comm_rank(MPI_COMM_WORLD, &myid);
    ...
    ierr = MPI_Finalize();
}
```



Advanced Research Computing
Enabling scientific breakthroughs at scale

MPI Example

```
#include
#include "mpi.h"

int main(int argc, char *argv[])
int argc;
char *argv[];
{
    int myid, nprocs;
    MPI_Init(&argc,&argv);
    MPI_Comm_size(MPI_COMM_WORLD,&nprocs);
    MPI_Comm_rank(MPI_COMM_WORLD,&myid);
    /* print out my rank and this run's NPROCS size*/
    printf("Hello from %d\n",myid," of ",nprocs);
    MPI_Finalize();
}
```



MPI Example: Send & Receive

```
#include "mpi.h"
/*****
This is a simple send/receive program in MPI
*****/
int main(int argc, char *argv[])
{
    int myid, numprocs, tag, source, destination, count, buffer ;
    MPI_Status status;
    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &numprocs);
    MPI_Comm_rank(MPI_COMM_WORLD, &myid);
    tag=1234;
    source=0;
    destination=1;
    count=1;
    if(myid == source){
        buffer=5678;
        MPI_Send(&buffer, count, MPI_INT, destination, tag, MPI_COMM_WORLD);
        printf("processor %d sent %d\n", myid, buffer);
    }
    if(myid == destination){
        MPI_Recv(&buffer, count, MPI_INT, source, tag, MPI_COMM_WORLD, &status);
        printf("processor %d got %d\n", myid, buffer);
    }
    MPI_Finalize();
}
```



Running MPI Jobs on Discovery

Compilers and MPI Stacks

- Available compilers on Discovery

Language	MPI Command (GCC, Intel)	Standard Command (GCC, Intel)
C	<code>mpicc, mpiicc</code>	<code>gcc, icc</code>
C++	<code>mpicxx, mpiicpc</code>	<code>g++, icpc</code>
Fortran 77/90	<code>mpif77/mpif90, mpiifort</code>	<code>gfortran, ifort</code>

- Available MPI stacks on Discovery: [mpich](#), [mvapich2](#), [openmpi](#), [intel-mpi](#)
- Run MPI jobs using Slurm's *srun* – use this within a Slurm job submission script

```
srun --mpi=pmix_v2 -n $SLURM_NTASKS ./mpi_program.x
```

← openmpi

```
srun --mpi=pmi2 -n $SLURM_NTASKS ./mpi_program.x
```

← mpich, mvapich2



MPI Job Slurm Scripts

```
#!/bin/bash
```

```
#SBATCH --nodes=6
```

```
#SBATCH --ntasks=144
```

```
#SBATCH --cpus-per-task=1
```

```
#SBATCH --mem-per-cpu=3GB
```

```
#SBATCH --time=24:00:00
```

```
#SBATCH --constraint="xeon-4116"
```

```
#SBATCH --exclusive
```

```
#SBATCH --account=<account_id>
```

```
module purge
```

```
module load gcc/8.3.0
```

```
module load openmpi/4.0.2
```

```
module load pmix/3.1.3
```

```
ulimit -s unlimited
```

```
srunk --mpi=pmix_v2 -n $SLURM_NTASKS ./mpi_program.x
```

144/6 = 24 (minimum number of cores available on each node)

Total number of process = total number of cores required (cpu-per-task=1)
This is the value assigned to \$SLURM_NTASKS, \$SLURM_NPROCS

A single core will be assigned to a single process

24 x 3GB = 72GB (minimum memory size required per node)

Asking xeon-4116 nodes only: is this the right choice? – check w/ *sinfo*

Not allowing other jobs to run on the selected nodes

Your account ID

Default compiler & mpi-stack in 'usc' module group

\$SLURM_NTASKS = 144

pmix_v2 is required for the use of openmpi



Advanced Research Computing
Enabling scientific breakthroughs at scale

Hands-on session w/ examples